

Reading in text data. Creating and visualizing a dfm

8 May 2019

In this document we will go through the steps of going from raw texts to a document term matrix that can be analyzed.

Reading in data

Let's take a look at a set of UK prime minister speeches from the EUSpeech dataset.

NB: Use `setwd()` to set the working directory to the folder that contains English speeches in the file `speeches_uk.csv`. Read in the speeches as follows using the `read.csv()` function:

```
Sys.setlocale(locale = "en_US.UTF-8")

## [1] "en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8"

speeches <- read.csv(file = "speeches_uk.csv",
                     header = TRUE,
                     stringsAsFactors = FALSE,
                     sep = ",",
                     encoding = "UTF-8")
```

Let's take a look at the structure of this dataset:

```
str(speeches)

## 'data.frame':    787 obs. of  6 variables:
## $ X      : int  1 2 3 4 5 6 7 8 9 10 ...
## $ text   : chr  "<p>This European Council has focused on 3 issues - the UK renegotiation, migration
## $ title  : chr  "EU Council: PM press conference" "PM statement in Poland: 10 December 2015" "PM st
## $ date   : chr  "18-12-2015" "10-12-2015" "09-12-2015" "07-12-2015" ...
## $ country: chr  "Great Britain" "Great Britain" "Great Britain" "Great Britain" ...
## $ speaker: chr  "D. Cameron" "D. Cameron" "D. Cameron" "D. Cameron" ...
```

As you can see, the corpus contains 787 speeches and variables containing meta data like speaker, country, date, etc. Take a look at a few speeches. Let's do some very light cleaning on these speeches, using the `stringr` library.

```
library(stringr)

#remove html tags
speeches$text <- str_replace_all(speeches$text, "<.*?>", "")
#replace multiple white spaces with single white spaces
speeches$text <- str_replace_all(speeches$text, " ", " ")
```

Question: If these were a proper data analysis, what other steps would you take to further clean the data?

Our `speeches` object is currently a dataframe. To be able to apply functions in `quanteda` on this object it needs to recognize it as a corpus. For this, you use the `corpus()` function

```
library(quanteda)
corpus <- corpus(speeches)

#the ndoc function displays the number of documents in the corpus
ndoc(corpus)
```

```
## [1] 787
```

Metadata such as speaker, date, etc. are stored in a corpus object as docvars, and can be accessed like so (we'll use the `head()` function to limit the output):

```
#date  
head(docvars(corpus, "date"), 10)
```

```
## [1] "18-12-2015" "10-12-2015" "09-12-2015" "07-12-2015" "07-12-2015"  
## [6] "01-12-2015" "28-11-2015" "23-11-2015" "19-11-2015" "16-11-2015"
```

```
#use the unique() function to check the number of unique speakers  
unique(docvars(corpus, "speaker"))
```

```
## [1] "D. Cameron" "G. Brown" "T. Blair"
```

```
#use the table() function to check the number of speeches for each speaker  
table(docvars(corpus, "speaker"))
```

```
##  
## D. Cameron G. Brown T. Blair  
##      493      283       11
```

```
#subsetting a corpus is easy using the corpus_subset() function  
cameron.corpus <- corpus_subset(corpus, speaker == "D. Cameron")  
ndoc(cameron.corpus)
```

```
## [1] 493
```

Sometimes it can be useful to tokenize your corpus. You can do this using the `tokens()` function.

```
tokens.speech <- tokens(corpus)
```

NB: the `tokens()` function generates a list that contains token vectors for all documents in the corpus. A list is an R object that can contain vectors of various lengths and types and as such is different from a dataframe which can contain variables of multiple types but with equal length.

Let's turn the corpus into a dfm using the `dfm()` function in `quanteda`:

```
corpus.dfm <- dfm(corpus, stem = FALSE, remove=stopwords("english"), remove_punct=TRUE)
```

Quanteda makes it very easy to inspect a dfm. For example, the `topfeatures()` function displays the most occurring features:

```
topfeatures(corpus.dfm, 20)
```

```
## people can think world want make just  
## 10628 9617 8848 5443 5119 4967 4817  
## now also need one country new going  
## 4632 4512 4423 4342 4342 4305 4257  
## countries britain prime us minister work  
## 4077 4056 4045 3921 3898 3884
```

You can check the number of features in the dfm using the `dim()` function:

```
dim(corpus.dfm)
```

```
## [1] 787 44533
```

There are over 44,000 features in this dfm. Let's select those tokens that appear in at least 10 documents by using the `dfm_trim()` function

```
corpus.dfm = dfm_trim(corpus.dfm, min_docfreq = 10)
dim(corpus.dfm)
```

```
## [1] 787 6104
```


As you can see, this reduces the size of the dfm considerably. However, be mindful that applying such arbitrary cutoffs may remove meaningful features.

NB: Because most words don't occur in most documents, a dtm often contains many zeroes (sparse). Internally, **quanteda** stores the dfm in a sparse format, which means that the zeroes are not stored, so you can create a dtm of many documents and many words without running into memory problems.

Visualization in quanteda

Quanteda contains some very useful functions to plot your corpus in order get a feel for it. For example, it is easy to construct a wordcloud to see which features appear most often in your corpus.

```
textplot_wordcloud(corpus.dfm, max_words=100)
```



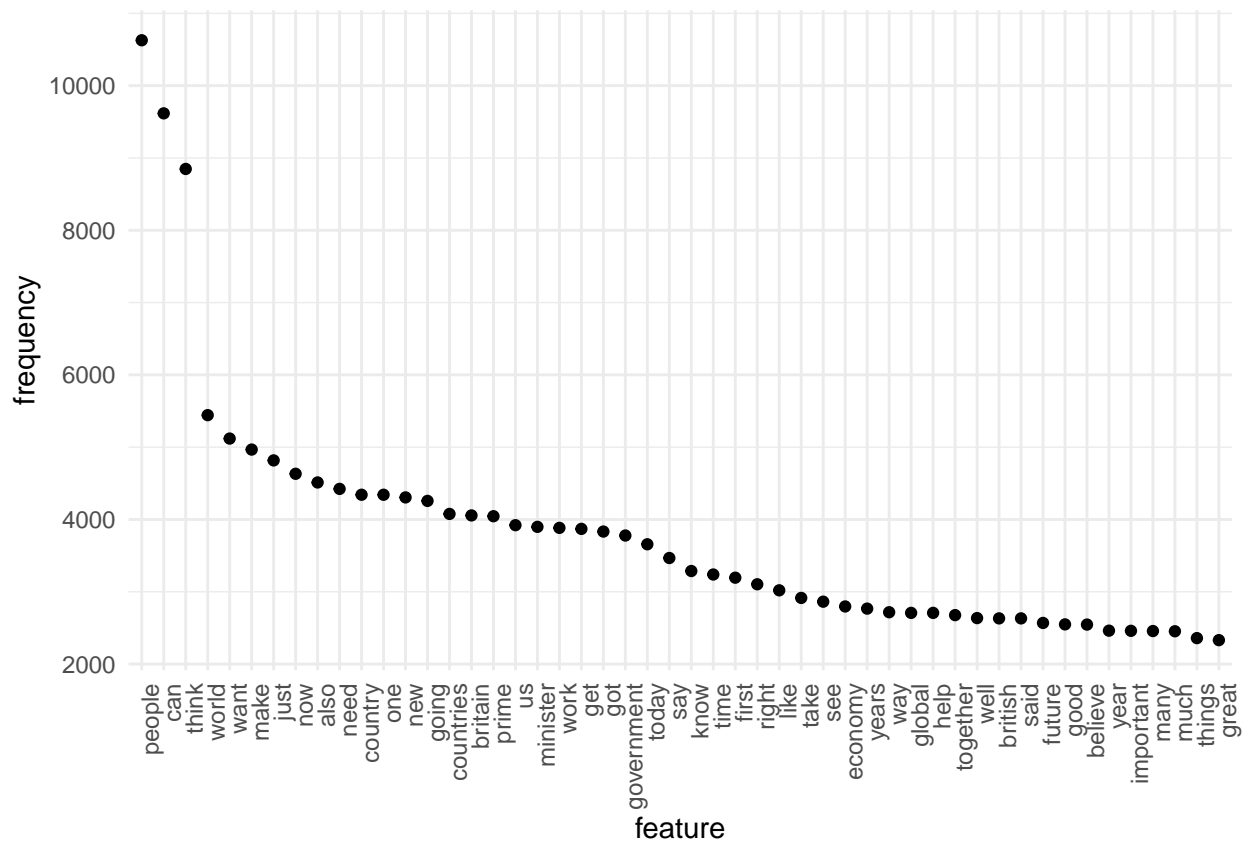
A more informative frequency plot can be constructed as follows (using the `ggplot2` library):

```
library(ggplot2)

corpus.dfm.features <- textstat_frequency(corpus.dfm, n = 50)

# Sort by reverse frequency order
corpus.dfm.features$feature <- with(corpus.dfm.features , reorder(feature, -frequency))

ggplot(corpus.dfm.features, aes(x = feature, y = frequency)) +
  geom_point() + theme_minimal() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1))
```




Let's say we are interested in which words are spoken relatively more often by David Cameron than by Tony Blair and Gordon Brown. For this we can use `textstat_keyness()` and `textplot_keyness()` functions.

```
cameron <- docvars(corpus, "speaker") == "D. Cameron"
comparison <- textstat_keyness(corpus.dfm, cameron)
head(comparison)
```

```
##   feature      chi2 p n_target n_reference
## 1    good 455.1106 0    1970      579
## 2     eu 333.2907 0     783      123
## 3 eurozone 331.5308 0     466       11
## 4   syria 302.1670 0     428       11
## 5   think 278.2528 0    5761    3087
## 6  things 247.5477 0    1708     651
```

```
textplot_keyness(comparison)
```



Question How would you interpret this plot?

Exercise

Take a look at this [Quanteda tutorial](#) for plotting frequency plots. Let's say you are interested in how often (in relative terms) David Cameron uses the word 'british' versus how often Gordon Brown uses it. Write some code to display this comparison in one plot.

NB The `ggplot2` is a really nice library for making plots and figures. If you have some time after this course is over, I strongly recommend Kieran Healy's book on Data Visualization for learning more about effective data viz.